

Regular Expressions in Python Tutorial

Michelle Fullwood

PyLadies Meetup Aug 2013

Motivating regular expressions

Scenario: You're evaluating Acme Company's products for a hospital. You're given a text file containing purchase records from Acme. Looking through the text file, however, you see that purchases from other companies are included as well, with no mention of which ones come from which!

PCOD	QTY	DEPT	COST
A169	100	Micro	0.58
PDA1	1	Xray	600.00
X280	5	ER	199.99
...			

Luckily, you know that Acme's product codes consist of one uppercase letter followed by three digits.

Motivating regular expressions

So you get to work with a script:

```
import string
for line in open('purchaserecords.txt', 'r'):
    if line[0] in string.uppercase and \
        line[1].isdigit() and \
        line[2].isdigit() and \
        line[3].isdigit():
        print line
    else:
        continue
```

It's a bit clunky, but it works.

Motivating regular expressions

The next step in your evaluation is to collate evaluations emailed to you by hospital staff. These were free text, so you need to extract the product codes from within them to know which evaluation refers to which product.

```
'...The gloves(P180) felt sticky...'  
'...The X701 vacuum cleaner really sucked!...'
```

You might be able to think of ways to program this, but really...

It's time to bust out regular expressions.

Motivating regular expressions

What we want is a way to simply search for “one uppercase letter followed by three digits”. We can do this using (1) a regular expression and (2) the `search` function provided by Python’s `re` module:

```
import re

re.search(r'[A-Z]\d{3}', mystring)
```

Just what are regular expressions, anyway?

Regular expressions are strings that describe other sets of strings.

$$\begin{matrix} \text{[A-Z]} & \text{\d{3}} \\ \left\{ \begin{matrix} A \\ B \\ \dots \\ Z \end{matrix} \right\} & \left\{ \begin{matrix} 0 \\ 1 \\ \dots \\ 9 \end{matrix} \right\} \end{matrix} \times 3$$

What we need to know to write regular expressions:

- ▶ How to define sets of characters
 - ▶ Metacharacters `\w`, `\s`, `\d`
 - ▶ Character sets `[A-Z]`, `[AGCT]`, `[^AGCT]`
- ▶ How to define how many times to repeat them
 - ▶ A specific number of times `{3,5}`, `?`
 - ▶ An unlimited number of times `*`, `+`

Plan for today

We'll learn:

- ▶ The pattern language for regular expressions
- ▶ The Python `re` functions that allow us to work with regexes

Playing along

How to practise the code as we go along:

- ▶ Install `git` (if you attended last month's meeting, you already have it!)
- ▶ From the command line:
`git clone https://github.com/michelleful/RegexTutorial.git`
- ▶ Run the exercises from the command line interface:
 - ▶ `cd exercises`
 - ▶ Open up `ex1a.py` etc as we move along
 - ▶ Edit these and run `python ex1a.py` from the CLI
- ▶ OR if you have iPython:
 - ▶ Run `ipython notebook` from the command line
 - ▶ A browser window will automatically open
 - ▶ Select the only notebook, follow the instructions therein
- ▶ OR if you don't have git, play with the regexes online
 - ▶ Enter regexes and strings into `http://www.pythonregex.com/`

Metacharacters

Metacharacters are pre-defined sets of characters.

- ▶ `.` matches ANY character*
- ▶ `\w` matches alphanumeric characters and underscore `_`
- ▶ `\d` matches digits 0 through 9
- ▶ `\s` matches whitespace characters
 - ▶ Spaces
 - ▶ Tabs `\t`
 - ▶ Newlines `\n\r`
 - ▶ (Escape sequences `\f\v`)

→ Exercise 1

Defining sets of characters

- ▶ List characters individually
 - ▶ `[AGCT]` matches one character A, G, C or T.
 - ▶ `[sd]` matches one whitespace character or digit
- ▶ Define a range of characters
 - ▶ `[A-T]` matches one character between A and T.
 - ▶ `[1-7]` matches one digit between 1 and 7.
 - ▶ Ranges as defined by ASCII or Unicode tables
 - ▶ You can combine ranges: `[a-cA-C]`

→ Exercise 2

Defining where a regular expression applies

We can say a regex has to be at the start or the end of the string, or at word boundaries, with more special characters.

- ▶ `^` – beginning of line
- ▶ `$` – end of line
- ▶ `\b` – word boundary

Examples: → Exercise 3

- ▶ `^Hallo$`
- ▶ `\bHallo\b`

Escaping characters

→ Exercise 4

Why didn't the `regex(es)` work? (Discuss.)

- ▶ When using characters that also have a special meaning, we have to escape them with a backslash
- ▶ `\^ \$ \. \\`
- ▶ Exception: within character sets `[]`, metacharacters have their regular meaning, except backslashes.*

Defining complements of a set

Sometimes it's easier to define a set of characters as “everything other than X”.

- ▶ `\S` – all non-whitespace characters
- ▶ `\W` – all non-alphanumeric characters (also excludes underscore)
- ▶ `\D` – all non-numeric characters
- ▶ `[^A-D]` – all characters other than A, B, C, D

→ Exercise 5

Repeating things

So far, whenever we've wanted to match, say, three digits, we've just been writing `\d\d\d`. This could easily get old if we want to match thirty or a thousand digits. It gets impossible when we want to match any arbitrary number of digits. Fortunately, regular expressions let us express this very succinctly.

Repeating things

- ▶ `*` means 0 or more times
- ▶ `+` means 1 or more times
- ▶ `?` means 0 or 1 times
- ▶ `{n}` means n times exactly
- ▶ `{m,}` means m or more times
- ▶ `{m,n}` means m–n times
- ▶ Example: `.*` matches anything
- ▶ Example: `[a-c]{3}` matches 'abc' etc

→ Exercise 6

Specifying alternatives

Sometimes you want to say match this OR that. You can do that with the `—` operator.

- ▶ `Alex|Bill|Conrad` matches any of these three names

Sometimes there can be confusion about what `—` refers to. In such cases, put brackets around the alternatives.

- ▶ `Jim and (Alex|Bill|Conrad)` matches 'Jim and Alex', 'Jim and Bill', etc

→ Exercise 7

Taking a step back

We've gone through a lot of the syntax for regular expressions.

1. Any questions?
2. Any challenging patterns you'd like to try as a group to express as a regex?

Python functions for dealing with regular expressions

- ▶ `re.find(regex, string)` – returns a match group if the string starts with the regex pattern, `None` otherwise.
- ▶ `re.search(regex, string)` – returns a match group if the string contains the regex pattern, `None` otherwise.

→ Exercise 8

Doing stuff with regular expressions

So far we've just been checking for matches with regexes. But we can do more!

We can:

- ▶ Capture the matches and do stuff with them
- ▶ Replace the match with something else
- ▶ Split a string whenever you match the regex (not covered today)

Capturing matches

We capture matches by putting brackets around the part we want to extract. Then access the matched bits via the match object returned by `.find()` and `.search()`

Example: We want to match strings that look like a first name followed by a last name.

Non-capturing parentheses

Replacing matches

Back references

Advanced topics

- ▶ Compiling regular expressions
- ▶ Flags
- ▶ Named matches

Compiling regular expressions

Sometimes you might see code like this:

```
import re

product_code_regex = re.compile(r'[A-Z]\d{3}')
print type(product_code_regex)
# <type '_sre.SRE_Pattern'>

print product_code_regex.findall('A103 B296 Z999')
# ['A103', 'B296', 'Z999']

# This is equivalent to:
print re.findall(r'[A-Z]\d{3}', 'A103 B296 Z999')
# ['A103', 'B296', 'Z999']
```

Once you have compiled your regex, you can call all the `re` functions directly as methods of the compiled regex.

Compiling your regex is a good idea but not necessary for small programs because Python will compile and cache your last few regexes anyway.

Flags

Flags change the way metacharacters, character sets and spaces are interpreted.

Some of the flags:

- ▶ `re.IGNORECASE` or `re.I` – case-insensitive matching
- ▶ `re.MULTILINE` or `re.M` – makes `^` and `$` match beginning and end of each newline
- ▶ `re.DOTALL` or `re.S` – makes `.` match any character *including* `\n`
- ▶ `re.UNICODE` or `re.U` – `\w`, `\b` reinterpreted according to Unicode definition
- ▶ `re.VERBOSE` or `re.V` – ignores whitespace and lets you use comments within the regex

Flags

How to set flags:

On any `re` function, add the flags as the last parameter:

- ▶ `re.search(myregex, mystring, re.U)`
- ▶ `re.compile(myregex, re.MULTILINE | re.DOTALL)`

Named matches

It's possible to identify groups by names:

```
# code from the Python docs
>>> m = re.match(r"(?P<first_name>\w+)
                  (?P<last_name>\w+)",
                  "Malcolm Reynolds")
>>> m.group('first_name')
'Malcolm'
>>> m.group('last_name')
'Reynolds'
```

This can be helpful for keeping track of multiple matches.

A word of caution on your newfound power...

*Some people, when confronted with a problem, think
“I know, I’ll use regular expressions.”
Now they have two problems.*

– Jamie Zawinski

When not to use regular expressions

Regular expressions are powerful, but you shouldn't reach for them whenever you need some string matching.

In particular, don't use regular expressions when:

- ▶ A simple string operation will do, e.g. `'abc' in mystring`
- ▶ You're parsing something that uses matching brackets, e.g. HTML and XML.
 - ▶ Use the `BeautifulSoup` or `lxml` libraries instead!